

AI BASED DIABETES PREDICTION SYSTEM

Development Part-1

V.SUMAN CSE-3rd Year B-sec

Introduction:

To Building a complete AI-based diabetes prediction system involves multiple steps, including obtaining the dataset, preprocessing the data, and building a predictive model. Below, I'll provide a step-by-step guide with code examples for each stage of the project.

I will setup the project environment, here are the steps I will be follow:-

- Download Anaconda,we use this to ease the process of installing necessary libraries for our project.
- Then I will create a repository and create a conda environment along with the necessary dependencies – Pandas, NumPy, Matplotlib and Scikit
- Activate the environment and install Jupyter notebook – an IDE to run our project.
- In Jupyter notebook, I will create a new python file and began the coding process.

Load the Dataset:

Load your dataset into a Pandas DataFrame. You can typically find kaggle Diabetes datasets in CSV format, but you can adapt this code to other formats as needed

To predict Diabetes using Diabetes dataset

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: dataset = pd.read_csv('Documents/PS/diabetes.csv')
```

```
[4]: dataset.head()
```

```
[4]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[5]: dataset.shape
```

```
[5]: (768, 9)
```

```
[6]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null    int64
1   Glucose                768 non-null    int64
2   BloodPressure          768 non-null    int64
3   SkinThickness          768 non-null    int64
```

```

4   Insulin          768 non-null   int64
5   BMI              768 non-null   float64
6   DiabetesPedigreeFunction  768 non-null   float64
7   Age              768 non-null   int64
8   Outcome          768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

```

```
[7]: dataset.describe().T
```

```

[7]:
      count      mean      std      min      25%  \
Pregnancies    768.0    3.845052    3.369578    0.000    1.000000
Glucose        768.0   120.894531   31.972618    0.000   99.000000
BloodPressure  768.0    69.105469   19.355807    0.000   62.000000
SkinThickness  768.0    20.536458   15.952218    0.000    0.000000
Insulin        768.0    79.799479  115.244002    0.000    0.000000
BMI            768.0    31.992578    7.884160    0.000   27.300000
DiabetesPedigreeFunction  768.0    0.471876    0.331329    0.078    0.24375
Age            768.0    33.240885   11.760232   21.000   24.000000
Outcome        768.0    0.348958    0.476951    0.000    0.000000

      50%      75%      max
Pregnancies    3.0000    6.00000    17.00
Glucose       117.0000   140.25000   199.00
BloodPressure   72.0000    80.00000   122.00
SkinThickness   23.0000    32.00000    99.00
Insulin        30.5000   127.25000   846.00
BMI            32.0000    36.60000    67.10
DiabetesPedigreeFunction   0.3725    0.62625    2.42
Age           29.0000   41.00000   81.00
Outcome        0.0000    1.00000    1.00

```

```
[8]: dataset.isnull().sum()
```

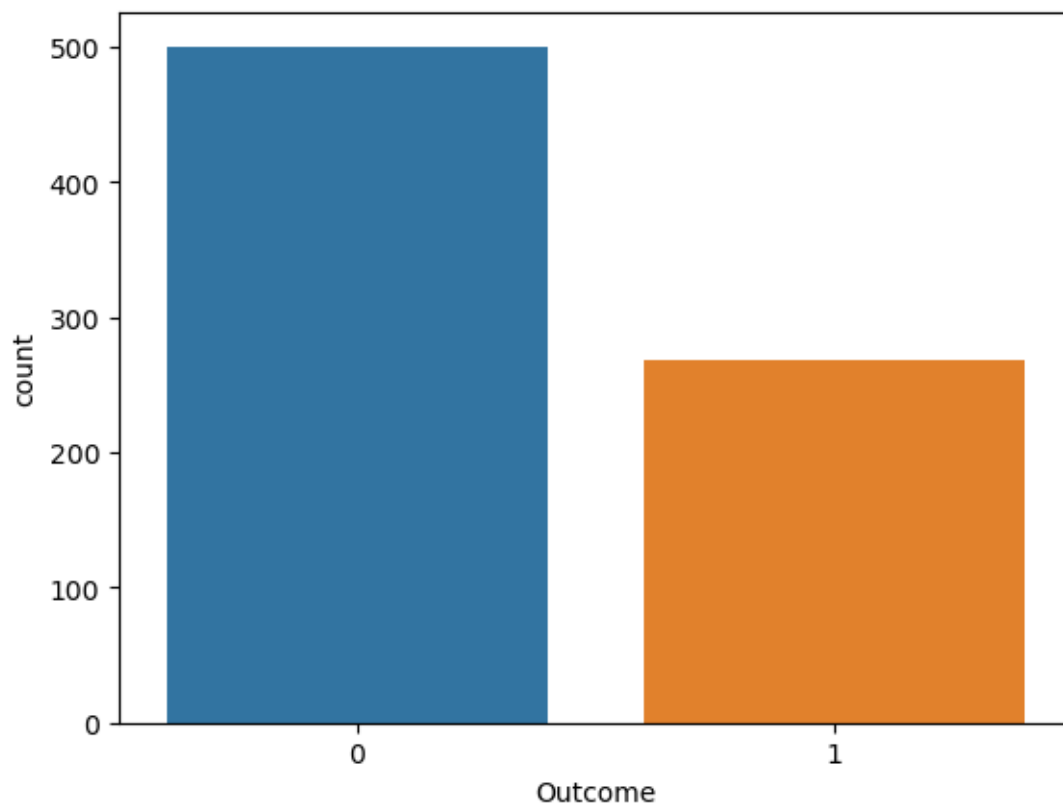
```

[8]: Pregnancies      0
      Glucose         0
      BloodPressure   0
      SkinThickness   0
      Insulin         0
      BMI             0
      DiabetesPedigreeFunction  0
      Age            0
      Outcome        0
dtype: int64

```

```
[9]: sns.countplot(x = 'Outcome',data = dataset)
```

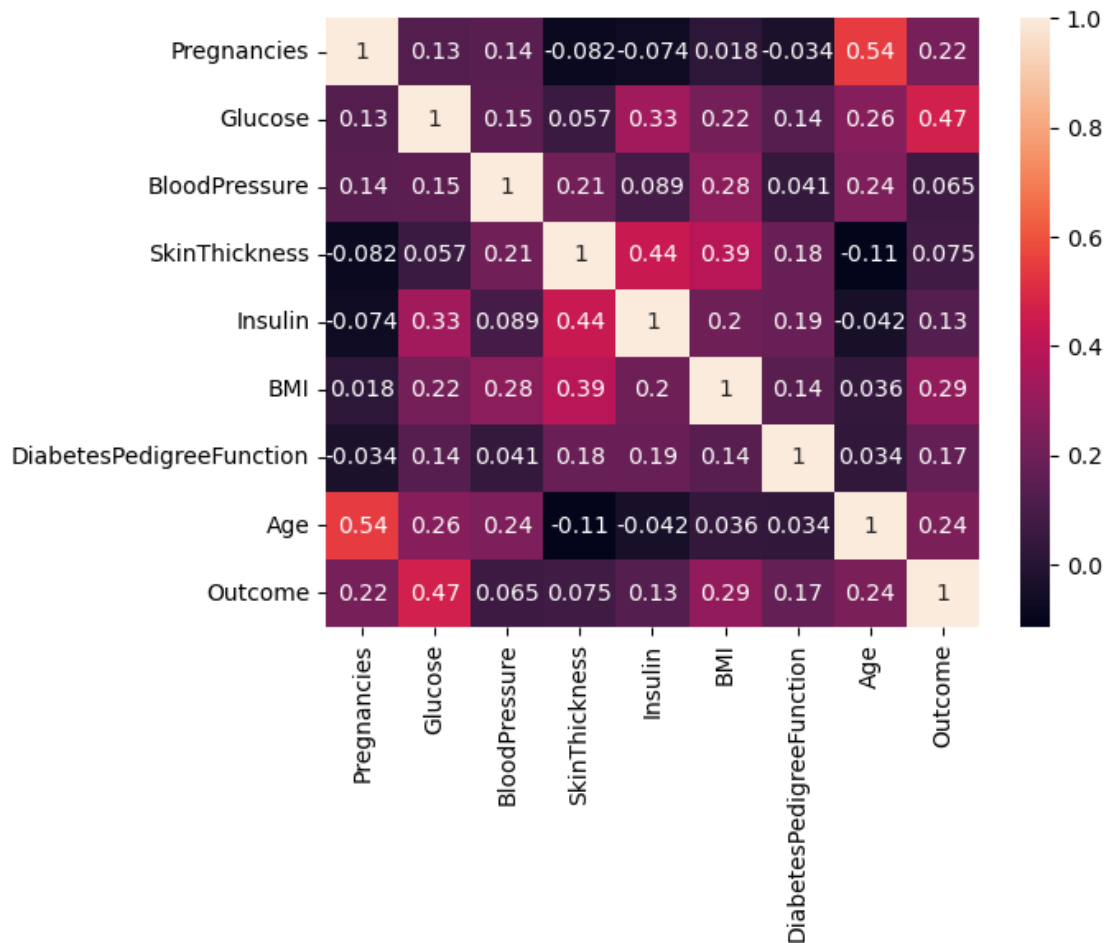
```
[9]: <Axes: xlabel='Outcome', ylabel='count'>
```



```
[12]: # Pairplot
sns.pairplot(data = dataset, hue = 'Outcome')
plt.show()
```



```
[13]: # Heatmap
sns.heatmap(dataset.corr(), annot = True)
plt.show()
```



```
[14]: # Replacing zero values with NaN
dataset_new = dataset
dataset_new[["Glucose", "BloodPressure", "SkinThickness", "Insulin", "BMI"]] =
    dataset_new[["Glucose", "BloodPressure", "SkinThickness", "Insulin", "BMI"]].
    replace(0, np.NaN)
```

```
[15]: # Count of NaN
dataset_new.isnull().sum()
```

```
[15]: Pregnancies      0
      Glucose         5
      BloodPressure   35
      SkinThickness   227
      Insulin         374
      BMI             11
      DiabetesPedigreeFunction  0
      Age             0
```

```
Outcome                                0
dtype: int64
```

```
[16]: # Replacing NaN with mean values
dataset_new["Glucose"].fillna(dataset_new["Glucose"].mean(), inplace = True)
dataset_new["BloodPressure"].fillna(dataset_new["BloodPressure"].mean(),
    ↪inplace = True)
dataset_new["SkinThickness"].fillna(dataset_new["SkinThickness"].mean(),
    ↪inplace = True)
dataset_new["Insulin"].fillna(dataset_new["Insulin"].mean(), inplace = True)
dataset_new["BMI"].fillna(dataset_new["BMI"].mean(), inplace = True)
```

```
[17]: dataset_new.isnull().sum()
```

```
[17]: Pregnancies                0
      Glucose                  0
      BloodPressure            0
      SkinThickness            0
      Insulin                  0
      BMI                      0
      DiabetesPedigreeFunction  0
      Age                      0
      Outcome                  0
dtype: int64
```

```
[18]: y = dataset_new['Outcome']
      X = dataset_new.drop('Outcome', axis=1)
```

```
[19]: # Splitting X and Y
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, y, test_size = 0.20,
    ↪random_state = 42, stratify = dataset_new['Outcome'] )
```

```
[20]: from sklearn.linear_model import LogisticRegression
      model = LogisticRegression()
      model.fit(X_train, Y_train)
      y_predict = model.predict(X_test)
```

```
C:\Users\CSE_BAY4\anaconda3\Lib\site-
packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-

```
regression
n_iter_i = _check_optimize_result(
```

```
[21]: y_predict
```

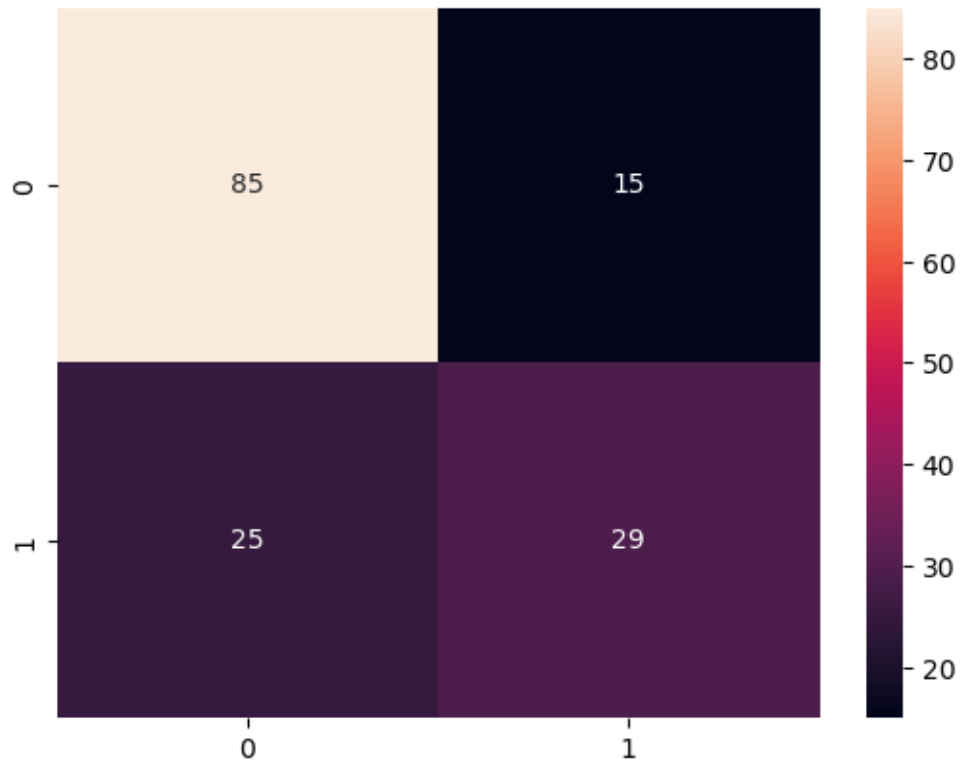
```
[21]: array([1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1,
          0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0,
          0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0,
          1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
          0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1,
          0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0,
          0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0],
          dtype=int64)
```

```
[22]: # Confusion matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(Y_test, y_predict)
cm
```

```
[22]: array([[85, 15],
          [25, 29]], dtype=int64)
```

```
[23]: # Heatmap of Confusion matrix
sns.heatmap(pd.DataFrame(cm), annot=True)
```

```
[23]: <Axes: >
```

```
[24]: from sklearn.metrics import accuracy_score
```

```
[25]: accuracy = accuracy_score(Y_test, y_predict)
accuracy
```

```
[25]: 0.7402597402597403
```

```
[26]: y_predict = model.predict([[1,148,72,35,79.799,33.6,0.627,50]])
print(y_predict)
if y_predict==1:
    print("Diabetic")
else:
    print("Non Diabetic")
```

```
[1]
Diabetic
```

```
C:\Users\CSE_BAY4\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning:
X does not have valid feature names, but LogisticRegression was fitted with
feature names
  warnings.warn(
```